

Pretrained Model

1.Object Detection

Pretrained Faster R-CNN with PyTorch (from TorchVision)

Logic Explained:

- **Model:** A two-stage detector (**Faster R-CNN**) pretrained on COCO (80 common classes).
- **Stage 1 – Region Proposal Network (RPN):** Slides over feature maps to propose candidate bounding boxes with "objectness" scores.
- **Stage 2 – Box classification & refinement:** Each proposed region is classified and its box coordinates fine-tuned.
- **Input pipeline:**
 1. Load image → convert to tensor $[C \times H \times W]$, batch dim $[1 \times C \times H \times W]$.
 2. Feed into model: `model = torchvision.models.detection.fasterrcnn_resnet50_fpn(weights="DEFAULT")`.
 3. Output: a dict with `boxes`, `labels`, `scores`.
- **Usage:** Loop through outputs, filter by score threshold (e.g. 0.5), then draw bounding boxes accordingly.

Link:

https://docs.pytorch.org/tutorials/intermediate/torchvision_tutorial.html?utm_source=chatgpt.com

2.Text Recognition (OCR)

EasyOCR for a simple OCR pipeline (*Adrian Rosebrock, PyImageSearch*)

Logic Explained:

- **Detection + Recognition pipeline all-in-one.**
 - **Initialization:** `reader = easyocr.Reader(['en'], gpu=False)`
 - **Read image:** `results = reader.readtext('image.jpg')`
 - **Outputs:** List of tuples (`bbox`, `text`, `confidence`)
 - `bbox`: 4 corner points of detected text region.
 - `text`: recognized string.
 - `confidence`: model's prediction score.
- **Internals breakdown:**
 - **CNN backbone** extracts features.
 - **LSTM** sequences spatial features.
 - **CTC decoding** maps sequences to characters.

Link: https://blog.roboflow.com/how-to-use-easyocr/?utm_source=chatgpt.com

3.Face Detection & Recognition

facenet-pytorch: MTCNN + FaceNet (InceptionResnetV1)

Logic Explained:

- **Face detection:**
 1. **MTCNN** – A 3-stage cascade (P-Net, R-Net, O-Net) to detect bounding boxes and landmarks.

2. Output: Locations of faces, plus cropped face tensor images.
- **Embedding generation:**
 1. `InceptionResnetV1(pretrained='vggface2').eval()` produces a 512-dim vector per face.
 - **Recognition:**
 1. Detect and align faces with MTCNN.
 2. Crop & normalize each face.
 3. Pass through the InceptionResnet model to get embeddings.
 4. Compare embeddings via cosine or euclidean distance.
 - If below a threshold (~0.8), they're the same person.
 - Uses FaceNet's triplet-loss training principle

Link:

https://www.reddit.com/r/pytorch/comments/xfcprn/tutorial_faster_rcnn_object_detection_with_pytorch/?utm_source=chatgpt.com